LET'S GO TO THE MOVIES!

Machine Learning & IMDB 5000 Movie Dataset

Team Corn on da CAWB
Claire Broad, Andre Duarte, Will Young, Brigit Lawrence-Gomez
MSAN621
December 8, 2016

Overview

Show me the money!

Question:
Can we predict the **total gross revenue** that a movie will make?

Techniques:
Random Forest
Decision Tree
KNN Regression

Ask yourself one question: 'Do I feel lucky?'

Question:
Can we **predict the genre** of a movie using keywords or movie description?

Techniques:
Logistic Regression
LDA

...the beginning of a beautiful friendship.

Question:
Can we find movies that are **similar to each other,** beyond looking at genre?

Technique:
Principal Component Analysis

The Data

- **IMDB 5000 data (https://www.kaggle.com/deepmatrix/imdb-5000-movie-dataset)**
  - This is the top 5000 US movies of all time, ranked by production budget

- **We scraped IMBD API for additional features:**
  - Release date
  - Plot description
  - Awards
  - Type (Movie or TV Series)
  - Writer
  - Metascore (a weighted critic score)

- **We also generated a few of our own features to help predict revenue:**
  - Average IMDB Rating score for each actor and director

Revenue Prediction: Feature Extraction

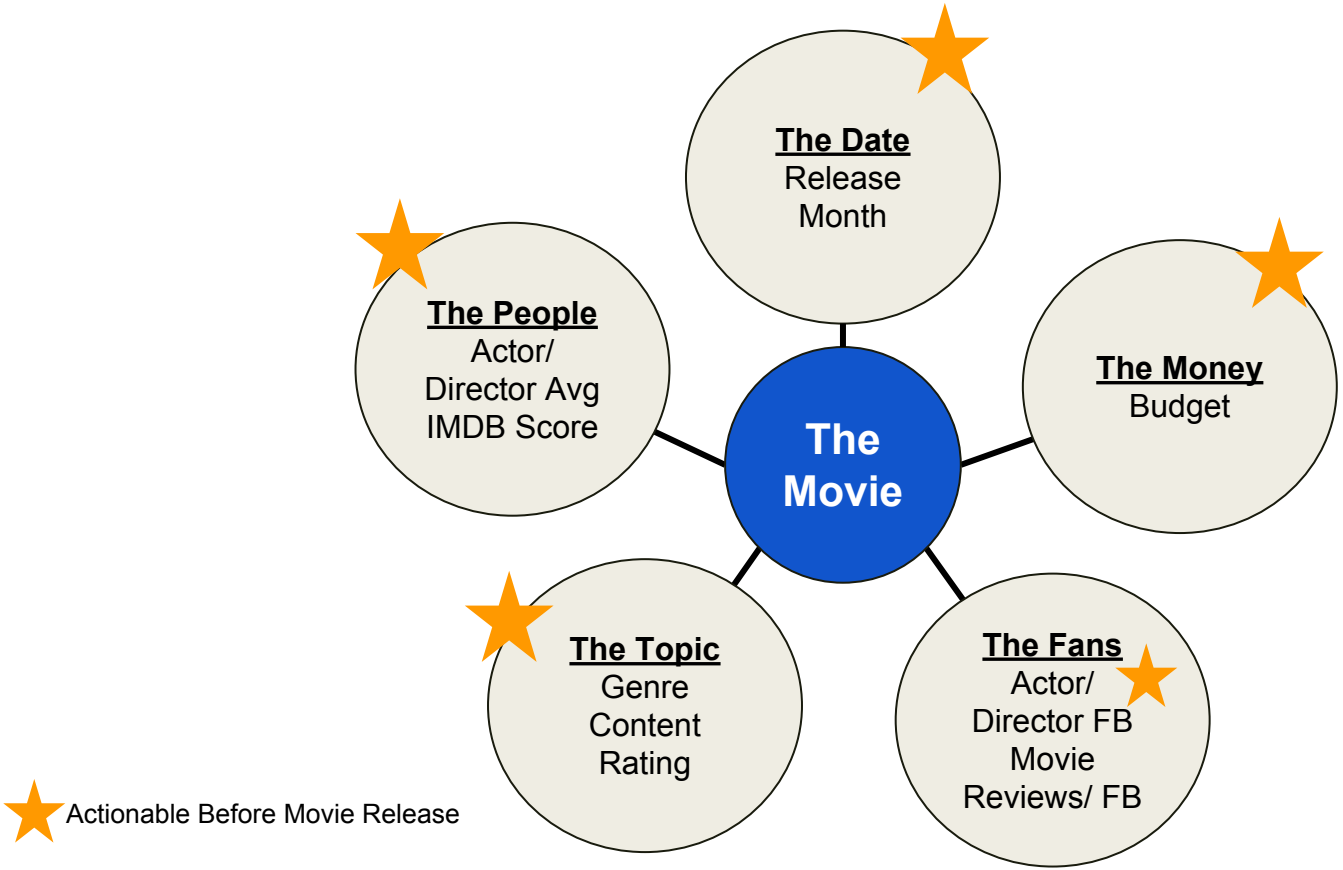

The Final Model: Random Forest Gump

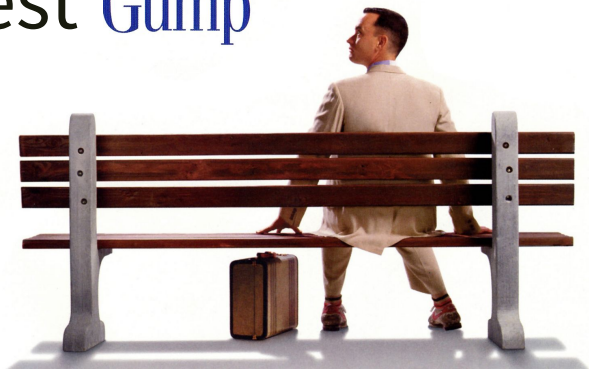

	All Features	Pre-Release Features Only
Tree Depth	12	28
Test RMSE	44e6	46e6
R^2	0.66	0.56
Gross Revenue Std. Dev. (Test)	76e6	69e6

Rank	All Features Model Feature	Pre-Release Only Model Feature
1	# Users Voted (IMDB)	Budget
2	Budget	Avg Director IMDB Score
3	# Users Wrote Reviews	Avg Lead Actor IMDB Score
4	R Content Rating	Duration
5	Avg Lead Actor IMDB Score	Avg 3rd Actor IMDB Score
6	Total Cast Facebook Likes	Total Cast Facebook Likes
7	Duration	Year
8	Year	Director Facebook Likes
9	Avg 2nd Actor IMDB Score	Avg 2nd Actor IMDB Score
10	Avg 3rd Actor IMDB Score	Release Month

Additional Insights

- **Content Rating** correlated with Revenue
  - Incentive for studio executives to pad rating
- **Total Budget** and Gross Revenue highly correlated and reveal seasonal trend
- **Total Marketing Budget** could potentially improve prediction to identify movies “hyped” in advance

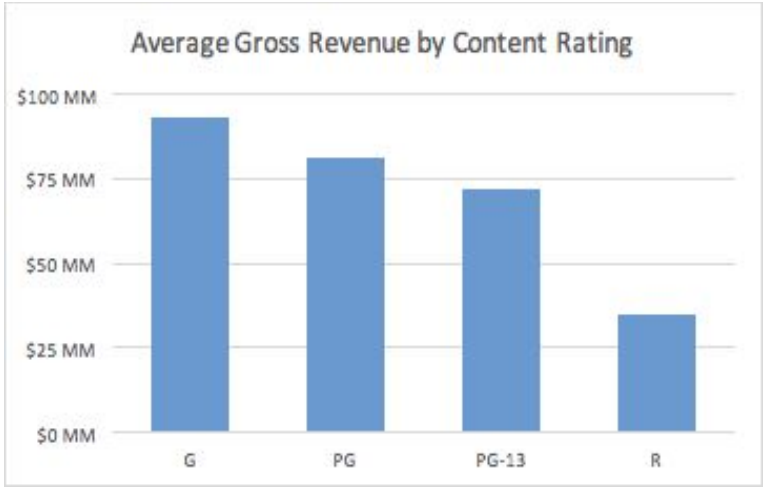

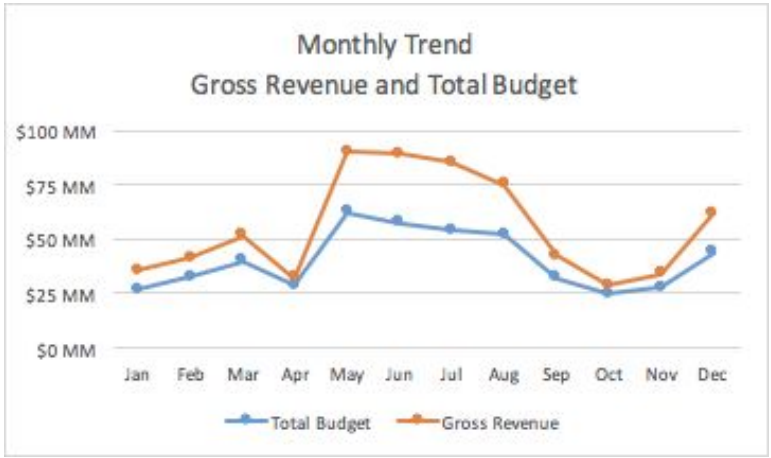

Genre Prediction

Given a 5 word description, can we guess the genre?

Action
Comedy
Drama
Romance
Fantasy
Horror
Thriller
True

Genre Prediction

Given a 5 word description, can we guess the genre?

Action
Comedy
Drama
Romance
Fantasy
Horror
Thriller
True

Genre Prediction

Given a 5 word description, can we guess the genre?

Correct: Action

Movie: Spiderman

Action
Comedy
Drama
Romance
Fantasy
Horror
Thriller
True

Genre Prediction

- Bucket genres into 8 categories
- Use top keywords as features
- Get accuracy score/ misclassification rate from several ML algorithms

Genre Prediction

- Bucket genres into 8 categories
- Use top keywords as features
- Get accuracy score/ misclassification rate from several ML algorithms

Problem: ‘Top 5’ keywords are not consistently good descriptors of the plot or tone

Genre Prediction

- Bucket genres into 8 categories
- Use top keywords as features
- Get accuracy score/ misclassification rate from several ML algorithms

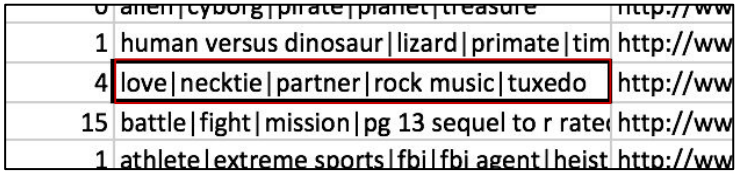

0	alien|cyborg|pirate|planet|treasure	http://ww
1	human versus dinosaur|lizard|primate|tim	http://ww
4	love|necktie|partner|rock music|tuxedo	http://ww
15	battle|fight|mission|pg 13 sequel to r rate	http://ww
1	athlete|extreme sports|fbi|fbi agent|heist	http://ww

Problem: 'Top 5' keywords are not consistently good descriptors of the plot or tone

Genre Prediction

- Bucket genres into 8 categories
- Use top keywords as features
- Get accuracy score/ misclassification rate from several ML algorithms

0	alien|cyborg|pirate|planet|treasure	http://ww
1	human versus dinosaur|lizard|primate|tim	http://ww
4	love|necktie|partner|rock music|tuxedo	http://ww
15	battle|fight|mission|pg 13 sequel to r rate	http://ww
1	athlete|extreme sports|fbi|fbi agent|heist	http://ww

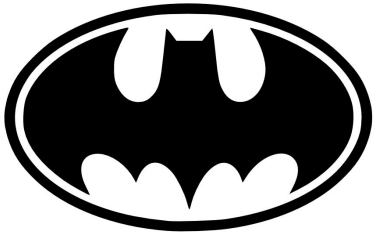

Problem: ‘Top 5’ keywords are not consistently good descriptors of the plot or tone

Genre Prediction

- Possible solution: Tokenized Plot Descriptions
- Pull top n word - genre pairs (n = 1000)
- Eliminate words with >4 genres
- 383 keywords remaining → features

```
Batman Forever
    Fantasy
        sidekick
        help            - Key Feature
        riddler
        young           - Key Feature
        batman
        circus
        two             - Key Feature
        face            - Key Feature
        psychologist
        becomes         - Key Feature
        battle          - Key Feature
        robin
        acrobat
        must            - Key Feature
```

Genre Prediction - Results

- Bucket accuracy
  - Keyword version: 40% (Logistic Regression)
  - Plot description version: 40% (LDA)

- Obscured accuracy
  - Keyword version: 20%
  - Plot description version: 17%

“Bad” Predictions

- Young Frankenstein
  - Keywords: assistant, castle, experiment, frankenstein's monster, scientist
  - Predicted: Horror
  - Actual: Comedy
- As Good as It Gets
  - Keywords: dog, friendship, neighbor, unlikely friendship, writer
  - Predicted: Comedy
  - Actual: Drama
- Batman Forever
  - Keywords: love, necktie, partner, rock music, tuxedo
  - Predicted: Drama
  - Actual: Fantasy

Good Predictions

- Pokémon 3: The Movie
  Keywords: ash, father, mother, pokemon, professor
  Bucket: Fantasy

- Friday the 13th: A New Beginning
  Keywords: jason voorhees, murder, new jersey, nightmare, teenager
  Bucket: Horror

- Transporter 2
  Keywords: driver, french, kidnapping, police, sequel
  Bucket: Thriller

PCA Clustering

- What are the most defining characteristics of movies in the dataset?
- What do the principal components mean?
- *Round up the **Usual Suspects**...*

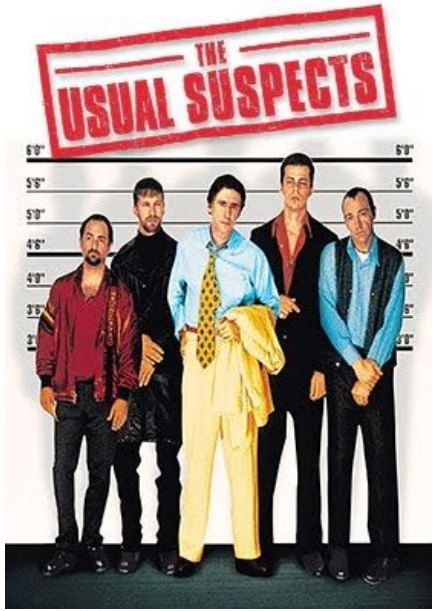

PCA Clustering

- **65%** of the Total Variance explained with **3** principal components

Percent of Total Variance
Explained by Principal Components

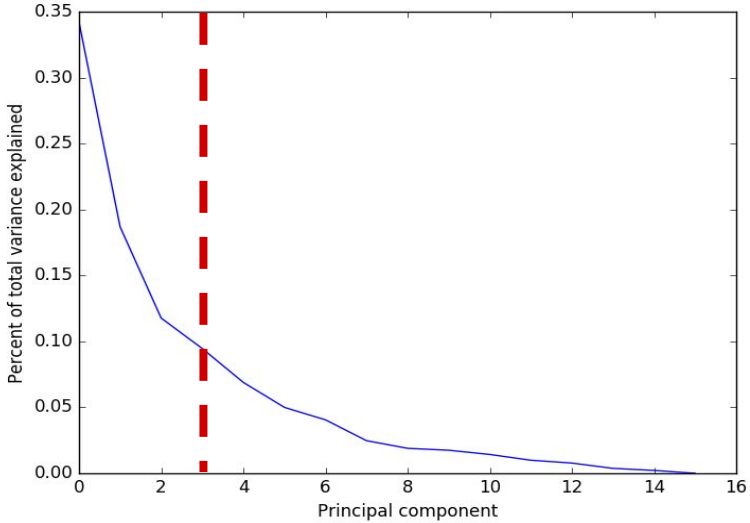

Cumulative Percent of Total Variance
Explained by Principal Components

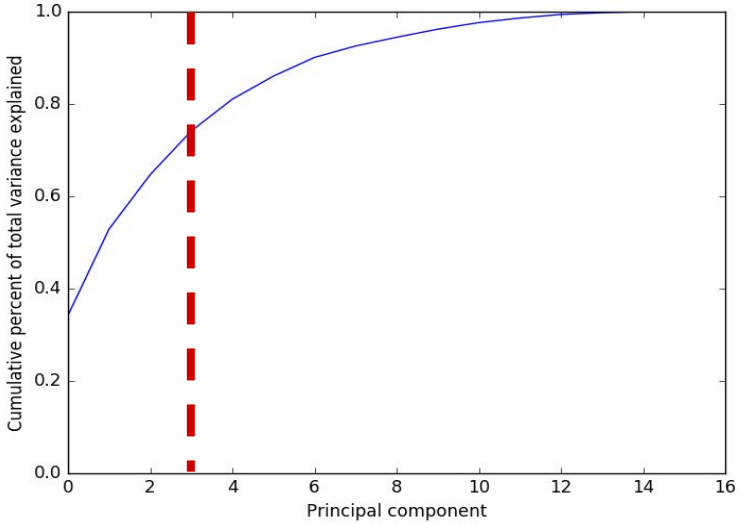

PCA Clustering

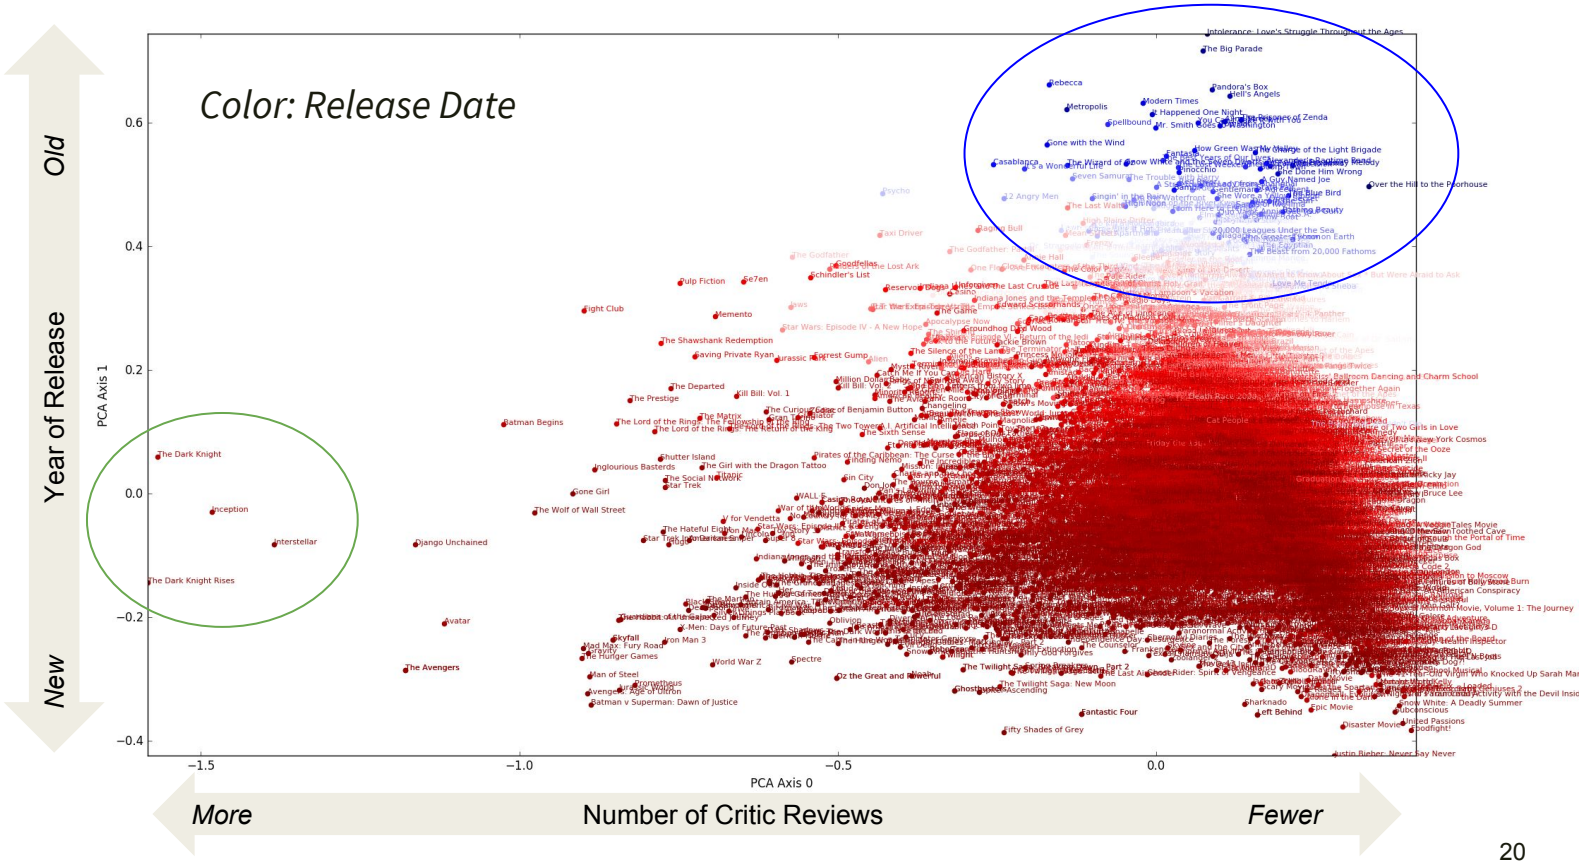


That's all Folks!


BLOOPERS

Revenue Prediction: Data Challenges

Data is like a box of chocolates. You never know what you're gonna get…

Data Challenges:

- Missing revenue data ~20% of the movies
- Movies combined with TV
- International movies had revenue in foreign currencies

Solution:

- Take them out

The Data

Movies in dataset skew more recent

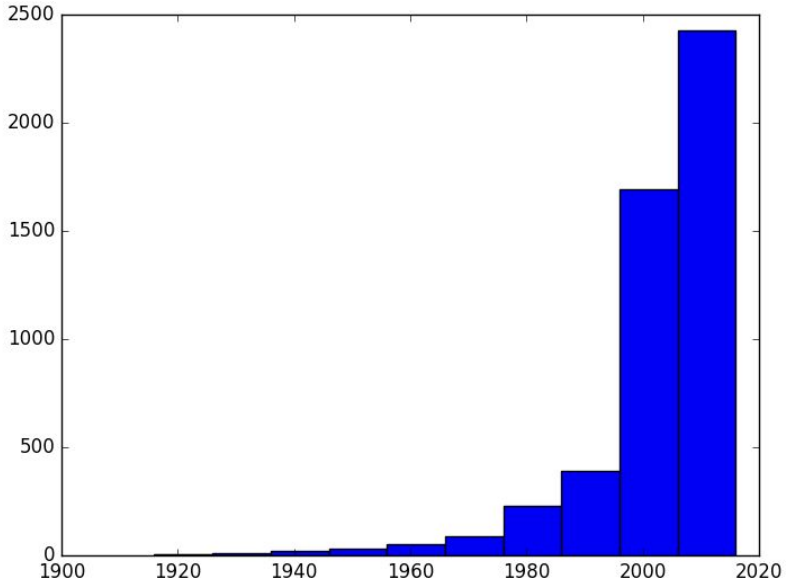

“Blockbusters” make revenue hard to predict